

# controlplane

Flux CD

## D1 Reference Architecture Guide

Rationale and Security considerations for the adoption of Flux CD D1 reference architecture on Kubernetes

# Table of Contents

Executive Summary	2
Introduction and Objective	3
Background	5
Organisation Suitability	5
User Personas	5
Low Level Design Description	8
Cluster Operating Model & Architecture	8
Example Workflows	12
Further Implementation Guidance	22
Securing GitHub	22
Multi-Tenancy Considerations	28
Other Security and Availability Considerations	33
About ControlPlane Flux Enterprise	35
Appendix	37
Administrative Tasks	37
About	44
Team	44
Reviewers	44

# Executive Summary

Flux CD, a cloud-native Continuous Deployment tool, has emerged as a crucial component in modern DevOps pipelines. In this guide, we delve into the intricacies of Flux CD, offering a comprehensive understanding of its functionality, implementation, and best practices.

Beginning with an insightful preface from the CNCF Flux maintainers, we embark on an exploration of Flux CD, starting with an introduction that sets the stage for the subsequent discussions. Our objective is clear: to equip readers with the knowledge and tools necessary to effectively leverage Flux CD within their organisations.

We provide essential background information, detailing the prerequisites and suitability of Flux CD for various organisational setups. Through the delineation of user personas and roles, we elucidate how Flux CD fits into different team dynamics, ensuring a tailored approach to implementation.

The guide delves into the low-level design description of Flux CD, elucidating the cluster operating model and architecture. We explore the intricacies of d1-fleet, d1-infra, and d1-apps repositories, offering a comprehensive understanding of their roles within the Flux ecosystem.

Drawing from real-world examples, we outline various workflows, including the process of adding a new application team and managing applications effectively. Furthermore, we offer implementation guidance, addressing crucial aspects such as GitHub as a production service, user access control, and authorization.

The document delves into Flux CLI usage and alternatives, namespace RBAC considerations, and addresses potential security concerns, including the tenant denial of service abuse case. It is also discussed at this point the role Kyverno policies are playing within Flux to help isolate the tenants.

Finally, readers will find additional information about ControlPlane Flux Enterprise, and an appendix containing the most common administration tasks.

This guide serves as a comprehensive guide for organisations seeking to harness the power of Flux CD for streamlined Continuous Deployment in cloud-native environments. This guide is mainly intended for those with some previous Flux knowledge but whether you are a seasoned DevOps practitioner or a newcomer to the field, this resource equips you with the knowledge and insights needed to succeed in your Flux CD journey.

# Introduction and Objective

In modern cloud native computing, efficiently and securely managing infrastructure and applications is crucial. Cloud-native technologies have exposed the limitations of traditional infrastructure management methods in meeting the demands of agility, scalability, security and reliability.

GitOps embodies the principles of Git, the popular version control system, to automate the deployment, monitoring, and management of cloud environments. GitOps is used for operating modern cloud native infrastructure and applications, at its core, GitOps treats infrastructure as code, enabling teams to define, track, and version their infrastructure configurations alongside application code.

GitOps promotes a declarative approach to infrastructure management, where the desired state of the system is defined in code and stored in a Git repository. Continuous synchronisation between the desired state described in the Git repository and the actual state of the system is achieved through automated workflows and reconciliation mechanisms.

As organisations continue to embrace cloud-native technologies and adopt DevOps practices, GitOps emerges as a foundational framework for managing cloud environments at scale. In this guide, we explore Flux, a leading GitOps tool, and its role in revolutionising cloud operations by streamlining the deployment, monitoring, and management of Kubernetes clusters and cloud-native applications. We delve into the features, benefits, and best practices of Flux, illustrating how it enables organisations to embrace GitOps principles and unlock the full potential of cloud-native ecosystems with improved security.

The purpose of this guide is to provide a comprehensive reference architecture demonstrating how Flux can address the needs of organisations with multiple teams looking to deploy applications on Kubernetes using GitOps principles. The guide aims to build upon the existing Flux documentation, showcasing how Flux can orchestrate multi-tenant clusters while catering to the different teams and stakeholders within an organisation.

This guide builds upon the existing Flux documentation, which features a quickstart guide orchestrating a single cluster, to demonstrate how Flux can orchestrate fleets of multi-tenant clusters.

It is also an important outcome of this guide to cater for the different teams and stakeholders that exist within an organisation (i.e. platform, services and application teams), whilst maintaining scalability, operational efficiency and security.

This guide is focused on the main challenges organisations face when adopting Flux CD in multi-team, multi-tenant environments. It provides solutions to difficulties such as:

- Segregating responsibilities and access between platform, services, and application teams
- Maintaining security and compliance in a multi-tenant environment
- Ensuring scalability and operational efficiency in large-scale deployments
- Implementing best practices for GitOps workflows and repository structures

After completing this guide, readers will have a clear understanding of how to implement a Flux CD reference architecture that caters to the needs of their specific organisational structure. They will be equipped with the knowledge and best practices to:

- Set up and manage multi-tenant Kubernetes clusters using Flux CD
- Define and maintain a repository structure that supports collaboration between platform, services, and application teams
- Implement secure GitOps workflows and access control mechanisms

The reference architecture detailed within the next sections can be deployed with either [Enterprise for Flux CD](#) or the open source version of Flux as both have feature parity. This document is the first of a series of reference architectures, designed to meet the needs of different organisational structures and ways of working, that will shortly be made available.

If you want to know the D1 Flux architecture and discover how to tackle the main security challenges it poses, bear with us for the rest of the document.

# Background

## Organisation Suitability

We have already walked you through the importance of GitOps, which challenges Flux enterprise architecture is aiming to solve and why there is a need for an enterprise version of Flux. The next points summarise the kind of organisation that would be suitable for this particular reference architecture implementation:

- Homogenisation of application packaging into containers
- Choice of Kubernetes as the orchestration platform
- A large number of clusters, or requirements to group applications resulting in multi-tenant, multi-cluster fleets
- *Namespace as a Service* operating model that involves a central platform team responsible for Kubernetes cluster provisioning, fleet maintenance and user experience with application teams as cluster tenants
- (Optional) Service teams responsible for maintaining and operating an individual service across a fleet of clusters (e.g observability/ security tools)
- Github qualifies as a production level service in terms of security and availability

## User Personas

When operating Kubernetes clusters at scale, it is critical to define the responsibilities of those which are setting up the Kubernetes environment, from those which are the users of that Kubernetes environment. Within this document we have named the former the platform team, meanwhile the latter is referred to as the application team.

In the next two sections you have access to a profile of these two roles.

## Platform Team

<i>Characteristic</i>	<i>Description</i>
Goals	<ul style="list-style-type: none"><li>• Provisioning and Maintenance of multi-tenant Kubernetes Clusters across the organisation</li><li>• Providing a tenant friendly experience</li><li>• Adhering to SLA Uptime and Security Requirements</li></ul>
Behaviours	<ul style="list-style-type: none"><li>• Planned Feature releases that are tested/ trialled in non-prod environments first. (These may include feature deprecations e.g Kubernetes API deprecations)</li><li>• Issues hotfix releases to all environments to address availability issues/bugs</li><li>• Will maintain the same configuration in all clusters</li><li>• Will use cluster admin console/CLI access to troubleshoot issues</li></ul>
Motivations	<ul style="list-style-type: none"><li>• Wants a simple way of managing clusters en-masse</li><li>• Would prefer minimal cluster specific configuration</li><li>• Want to eradicate cluster configuration drift</li></ul>
Pain Points	<ul style="list-style-type: none"><li>• Application teams not configuring their applications in accordance with best practice</li><li>• Managing releases at scale, obtaining a feedback loop to understand when new releases break application workloads</li></ul>
Risk Behaviours	<ul style="list-style-type: none"><li>• Will not apply security policy to own namespaces and will run privileged platform workloads without assessing for least privilege on the cluster</li></ul>

## Application Team

<i>Characteristic</i>	<i>Description</i>
Goals	<ul style="list-style-type: none"><li>• Building and running containerised applications on Kubernetes</li></ul>
Behaviours	<ul style="list-style-type: none"><li>• Will have different configuration between non-prod and production environments (e.g log levels, external endpoint addresses)</li><li>• Will require multiple namespaces in non-prod clusters</li><li>• Will aim to release as frequently as possible.</li></ul>
Motivations	<ul style="list-style-type: none"><li>• Simple mechanism of consuming Kubernetes</li><li>• Instant deployment feedback</li></ul>
Pain Points	<ul style="list-style-type: none"><li>• May not be Kubernetes configuration experts</li><li>• Kubernetes platform changes may cause breaking changes to their application</li></ul>
Risk Behaviours	<ul style="list-style-type: none"><li>• Will configure their applications by any means possible to get it working, which without preventive controls could cause security misconfigurations, or unusual configurations for the platform team to support.</li><li>• Will turn off namespace level controls to get their applications working if required</li></ul>

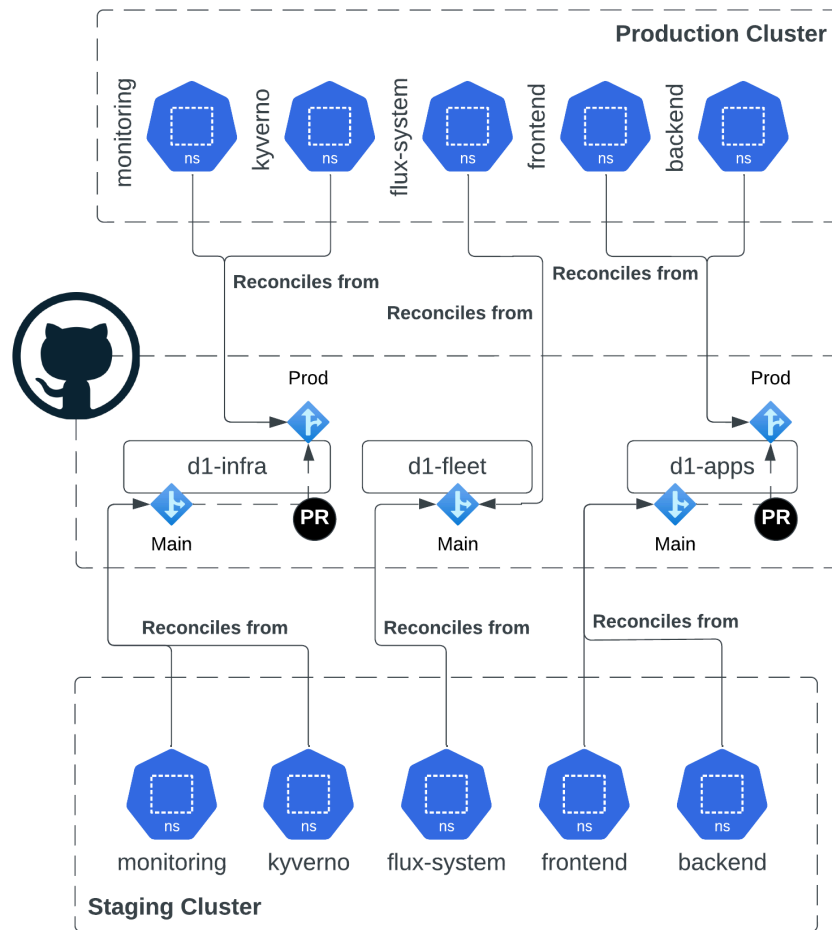


# Low Level Design Description

## Cluster Operating Model & Architecture

### D1 Architecture Repositories Structure

To enable the requirements and needs of the user personas described within the previous section, the next repository structure has been created:



**Figure 1: D1 Architecture - repositories structure.**

### D1-Fleet Repository

Utilised by the platform team only, who are admin on all clusters in order to:

- Bootstrap Flux with multi-tenancy restrictions on fleet clusters
- Configure the delivery of platform components (defined in d1-infra repository)
- Configure the delivery of applications (defined in d1-apps repository)

## D1-Infra Repository

This repository is managed by the platform team who are responsible for the Kubernetes infrastructure. This repository is used to define the Kubernetes infrastructure components such as:

- Cluster add-ons (CRD controllers, admission controllers, monitoring, logging, etc.)
- Cluster-wide definitions (Namespaces, Ingress classes, Storage classes, etc.)
- Pod security standards
- Network policies

In larger organisations, where dedicated teams may be responsible for services (e.g a security team managing policies, secrets team managing secrets operators, or observability team managing prometheus), these teams will be configuring their components within this repository.

## D1-Apps Repository

Each development team responsible for the delivery of an application running on the Kubernetes cluster fleet will be allocated an Application Repository, hosting application components such as:

- Flux HelmRepositories (pointing to the application Helm charts in container registries)
- Flux HelmReleases for the applications with custom configuration per environment

This repository is reconciled on the cluster fleet by Flux as the namespace admin and can't contain Kubernetes cluster-wide definitions such as CRDs, Cluster Roles, Namespaces, etc.

Access to this repository is restricted to the dev teams and the [Flux bot account](#). The Flux bot should be the only account with direct push access to the main branch.

## Branching Strategy

Flux instances, and Kustomizations/GitRepositories pointers to the d1-infra and d1-apps repositories on all clusters are configured off the main branch of the d1-fleet repo.

The infra components and applications on staging clusters are controlled by the main branch of the respective repositories, whilst production clusters are

synchronised against the production branch. This permits the Flux image automation controllers to automatically deploy updated Helm Charts in staging whilst ensuring that protected branches (a Pull Request is needed to merge the content into the production branch) and appropriate release controls are applied to the production environment. See [Protected branches](#) for further detail.

You can also refer to the diagram within the [D1 Architecture Repositories Structure](#) section of the document.

## Update Automation

Through image automation controllers, Flux can enable continuous deployment such that new application versions can be deployed to any environment. Within the D1 architecture, this is used to automate deployments to Staging clusters.

During the staging cluster bootstrap, the `image-reflector-controller` and `image-automation-controller` are deployed (see [here](#)). It is not necessary to deploy the automation components to production clusters. Pushes to the main branches will be performed by at least one staging cluster.

These controllers work together to automatically update the `HelmReleases` stored in the main branch. More details on the process can be found [here](#).

This automation is built around different components:

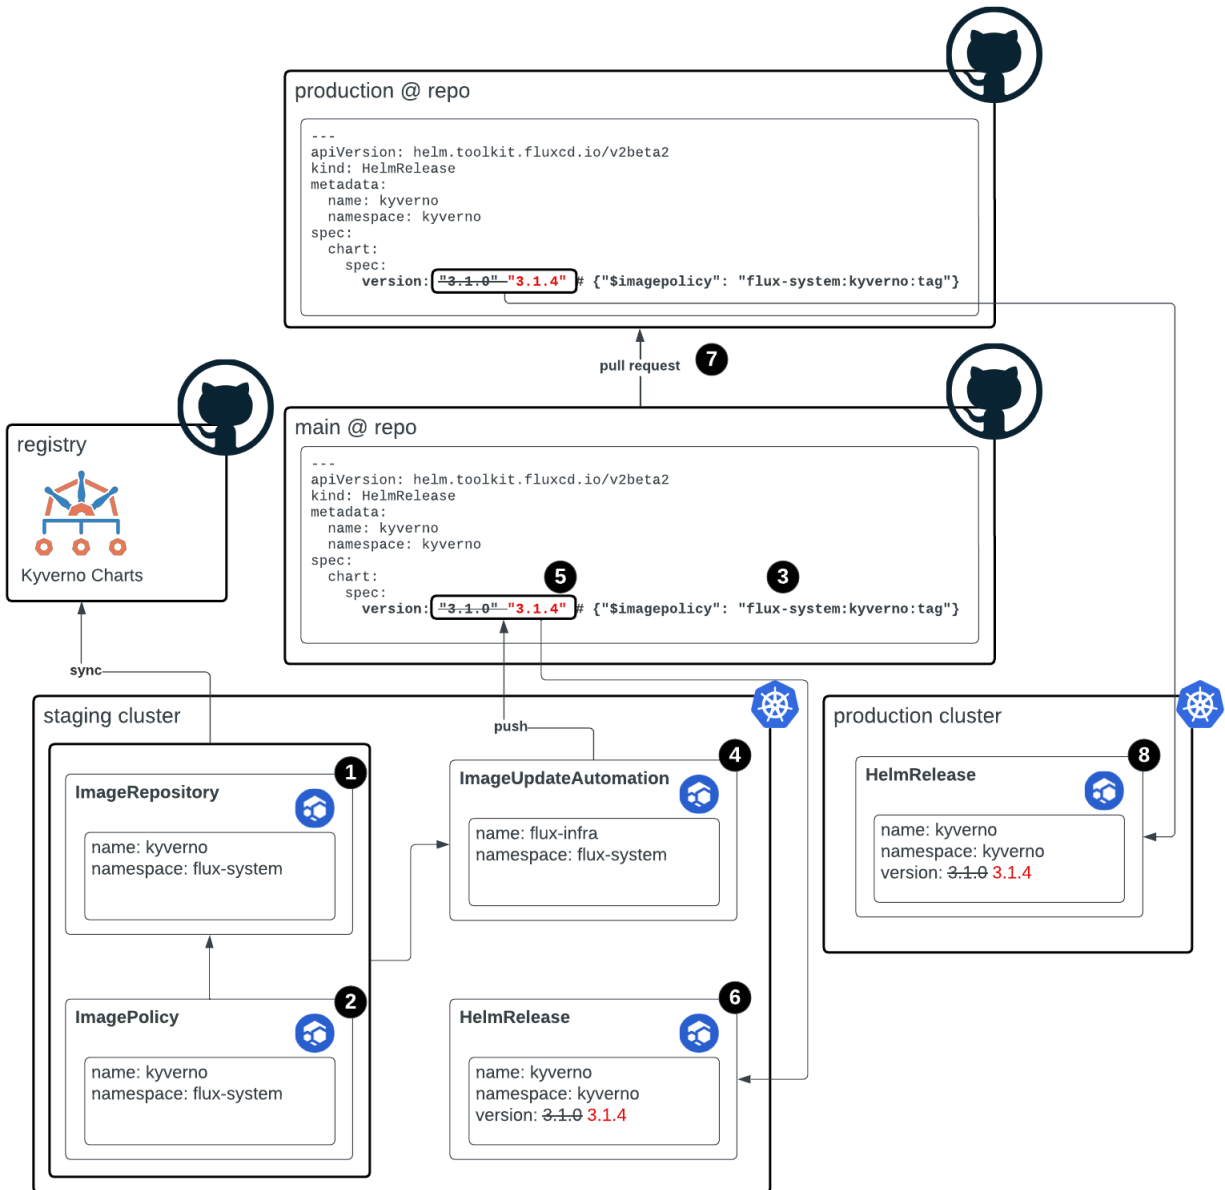
1. `ImageRepository` CR. This tells Flux which container registry to scan for new tags
2. `ImagePolicy` CR. This tells Flux which [semver](#) range to use when filtering tags
3. Specific markers on the affected manifests to tell where Flux should update the version and with which policy (i.e. in a `HelmRelease.spec.chart.spec.version`)
4. `ImageUpdateAutomation` CR. This tells Flux which Git repository to write image updates to. The PAT associated with the `GitRepository` CR used by `ImageUpdateAutomation` must have enough privileges to perform pushes on the main branch (see [here](#))

As soon as Flux will detect an update then:

5. The `HelmRelease` manifests will be updated on the main branch;
6. The `HelmRelease` on all the staging clusters will be reconciled to the new version.

To update production clusters:

7. Open and approve a pull request to the production branch;
8. Wait for Flux to reconcile the updated HelmReleases on all the production clusters.



**Figure 2: Update Automation.**

Following the dI model, **ImageRepository**, **ImagePolicy** and manifest markers are pushed to the repository by the tenant responsables. **ImageUpdateAutomation** CR and the RBAC needed to reconcile the update automation components are deployed on the fleet repository by the platform team for each tenant.

# Example Workflows

## Cluster Onboarding

The bootstrap procedure is a one-time operation that sets up the Flux controllers on the cluster (in this case a staging cluster), and configures the delivery of platform components and applications.

The bootstrap procedure requires the following:

- Flux CLI
- A Github Bot Account Personal Access Token (PAT) with [suitable permissions](#)
- Cluster admin permissions on a provisioned Kubernetes cluster

The Flux CLI or Terraform Module can be called as a step within a cluster provisioning pipeline or used from an administrator's machine. See [flux bootstrap command](#) for further details and further discussion on Flux CLI usage can be found in [Flux CLI usage & Alternatives](#)

The Flux CLI will use the bot PAT to push two commits to the fleet repository:

1. First commit to create the clusters/staging/flux-system/gotk-components.yaml file which contains the flux-system namespace, RBAC, network policies, CRDs and the controller deployments
2. Second commit to create the clusters/staging/flux-system/gotk-sync.yaml file which contains the Flux GitRepository and Kustomization custom resources for setting up the cluster reconciliation

This Flux CLI will perform the following actions on the cluster:

3. Creates a Kubernetes Secret named flux-system in the flux-system namespace that contains the bot PAT
4. Builds the cluster/staging/flux-system kustomize overlays with the multi-tenancy patches and applies the generated manifests to the cluster to kick off the reconciliation

From this point on, the Flux controllers will:

5. Reconcile the cluster state with the desired state within the fleet repository
6. The tenant folder within the fleet repository contains GitRepository CRs that point to d1-infra and d1-apps repositories as further sources, as well as bootstrap configuration (namespace/RBAC) for those tenants
7. The Flux controllers subsequently reconcile the cluster state with the desired state within the d1-infra and d1-apps repositories

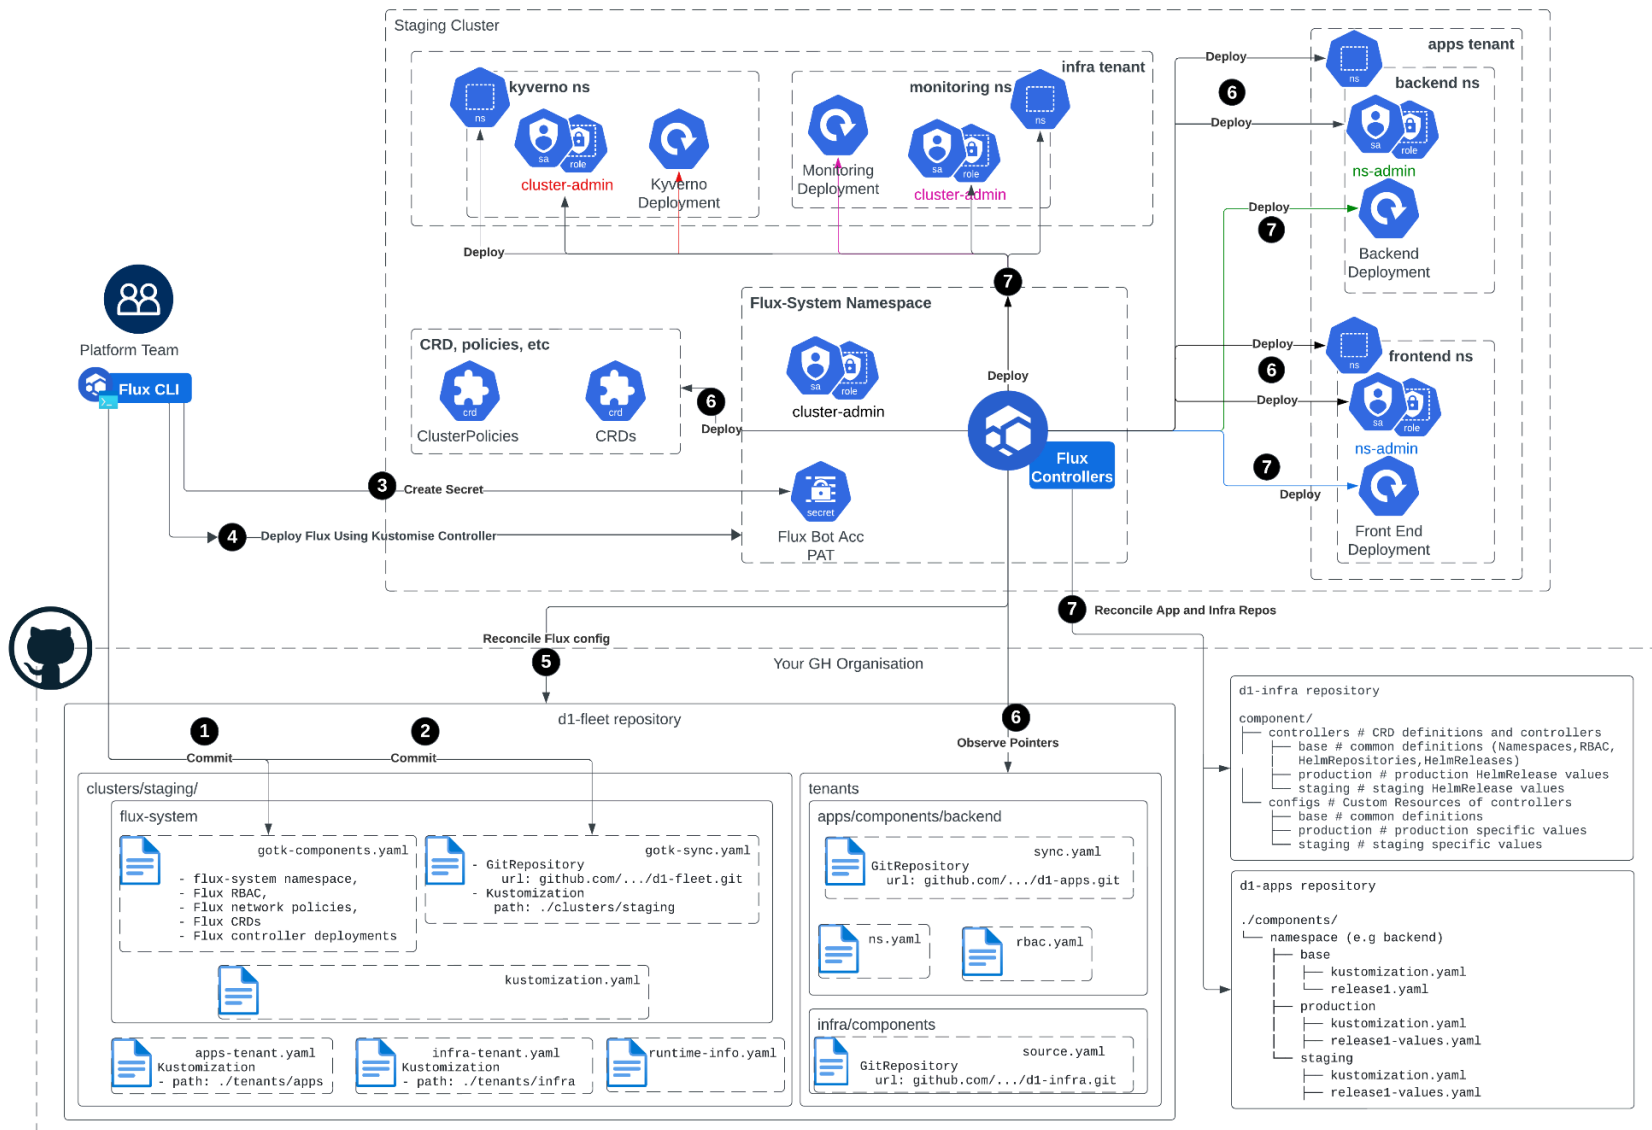


Figure 3: Example Workflow - Cluster onboarding.

## Adding a New Infra Component

This process is followed by the platform team to add a new infrastructural component to the fleet. The diagram below shows where and when the manifests for a new infra component are pushed to the different repositories.

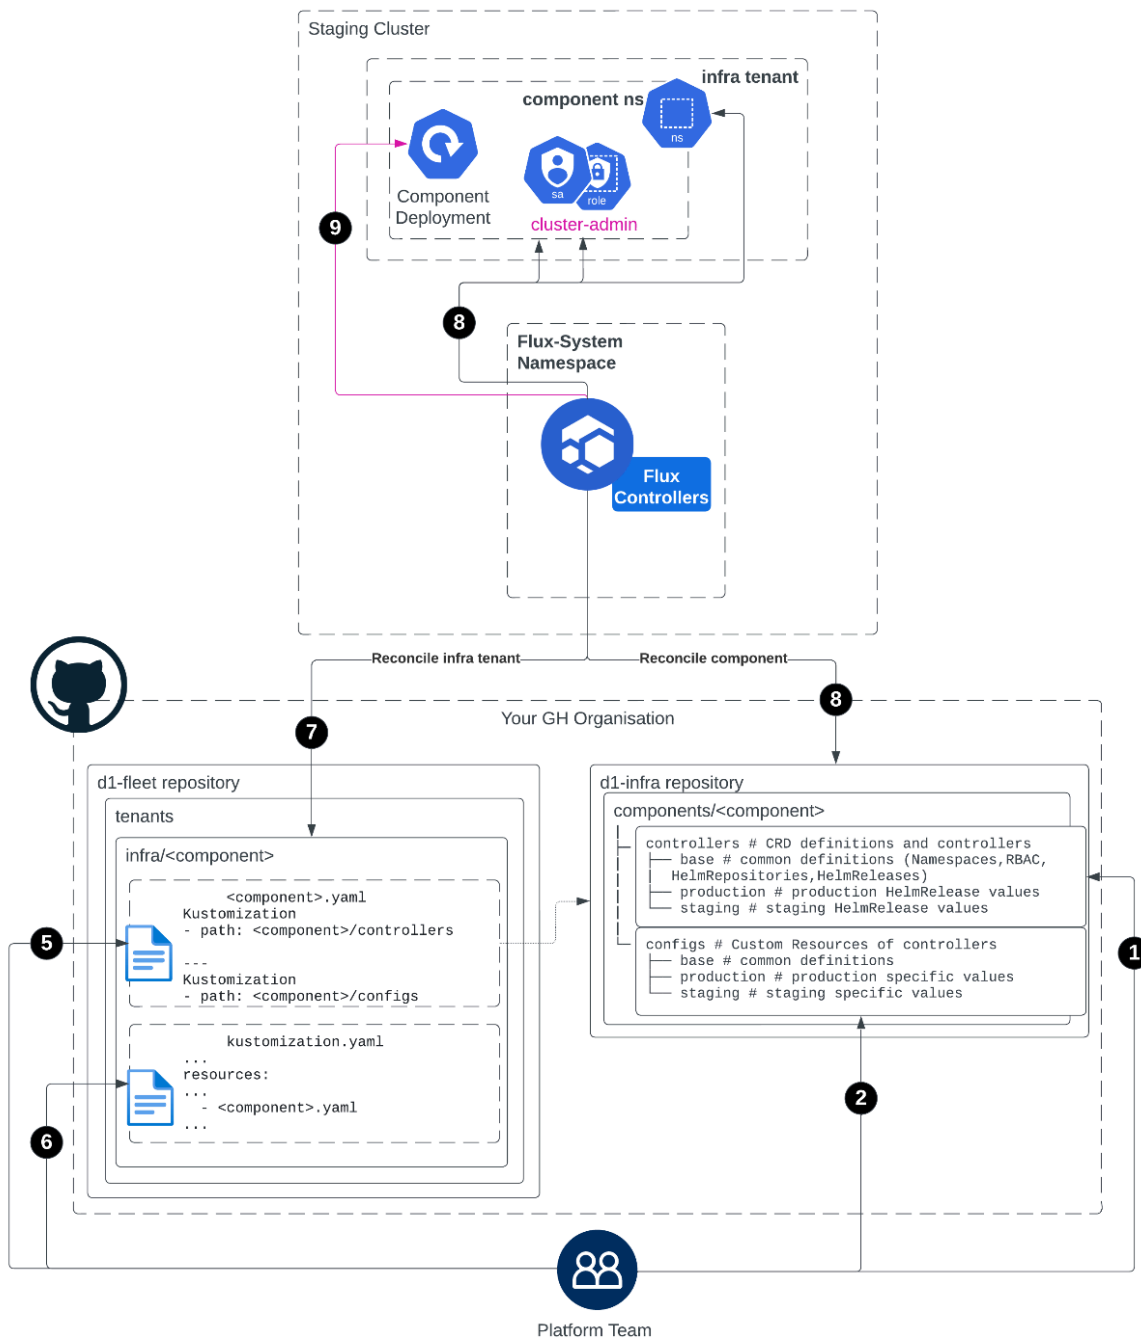


Figure 4: Example Workflow - Adding a new infrastructure component.

To onboard a component from the d1-infra repository, the platform team should fulfil the following requirement:

- Read-write access to d1-fleet repository
- Read-write access to d1-infra repository, on main and on production branch

Platform team should then push on d1-infra main branch:

1. Add the infra component manifests and overlay to the [d1-infra]/components/<component>/controllers folder, in order to
  - a. Deploy the needed **Namespace**
  - b. Deploy a new Flux **ServiceAccount** that is going to be responsible of the reconciliation for all the controllers belonging to that namespace
  - c. Deploy a new **ClusterRoleBinding** that will assign **cluster-admin ClusterRole** to the **Flux SA**
  - d. Deploy the component itself (i.e. using **HelmRelease**) and any needed CRDs
2. Add the needed manifests and overlay to the [d1-infra]/components/<component>/configs, in order to deploy any configuration/custom resources used by that component
3. (optional) automatize updates:
  - a. Add **ImageRepository** and **ImagePolicy** CR to [d1-infra]/update folder and reference it in [d1-infra]/kustomization.yaml manifest
  - b. Add all the needed markers for the update automation workflow

Then, the team should push on d1-fleet main branch to start the reconciliation:

4. (Optional) add any useful cluster variable to [d1-fleet]/clusters/<cluster>/runtime-info.yaml **ConfigMap**
5. Add new **Kustomizations** in [d1-fleet]/tenants/infra/components that will act as a pointer to the newly created manifests in **controllers** and **configs** folders of d1-infra repo. An explicit dependency between controllers and configs has to be declared, in order for the controller to be deployed before the configs
6. Add the reference to this **Kustomization** in [d1-fleet]/tenants/infra/components/kustomization.yaml in order to kickoff the reconciliation for this component

Flux will then:

7. Reconcile the infra tenant and then notice two new **Kustomizations** for the component (for **controllers** and **configs**)
8. Reconcile the RBAC and namespace from d1-infra/components/<component>/controllers folder (**HelmRepository**, **HelmRelease**, etc)



9. Reconcile the `HelmRepository`, `HelmRelease` from `d1-infra/components/<component>/controllers` folder
10. Reconcile the objects from `d1-infra/components/<component>/configs` folder

Finally, the platform team will open and approve a pull request on the `d1-infra` production branch to apply the changes in production.

An infrastructural tenant might be organised in the following way

```
Unset
component/
├─ controllers # CRD definitions and controllers
│   └─ base # common definitions (NS, RBAC, HelmRepositories, HelmReleases)
│   └─ production # production specific HelmRelease values
│   └─ staging # staging specific HelmRelease values
└─ configs # Custom Resources of controllers
    └─ base # common definitions
    └─ production # production specific values
    └─ staging # staging specific values
```

You can find a practical example in the [Appendix](#).

## Adding a New Application Team/Tenant

The platform team should be responsible for adding a new tenant for each application team, following the process below.

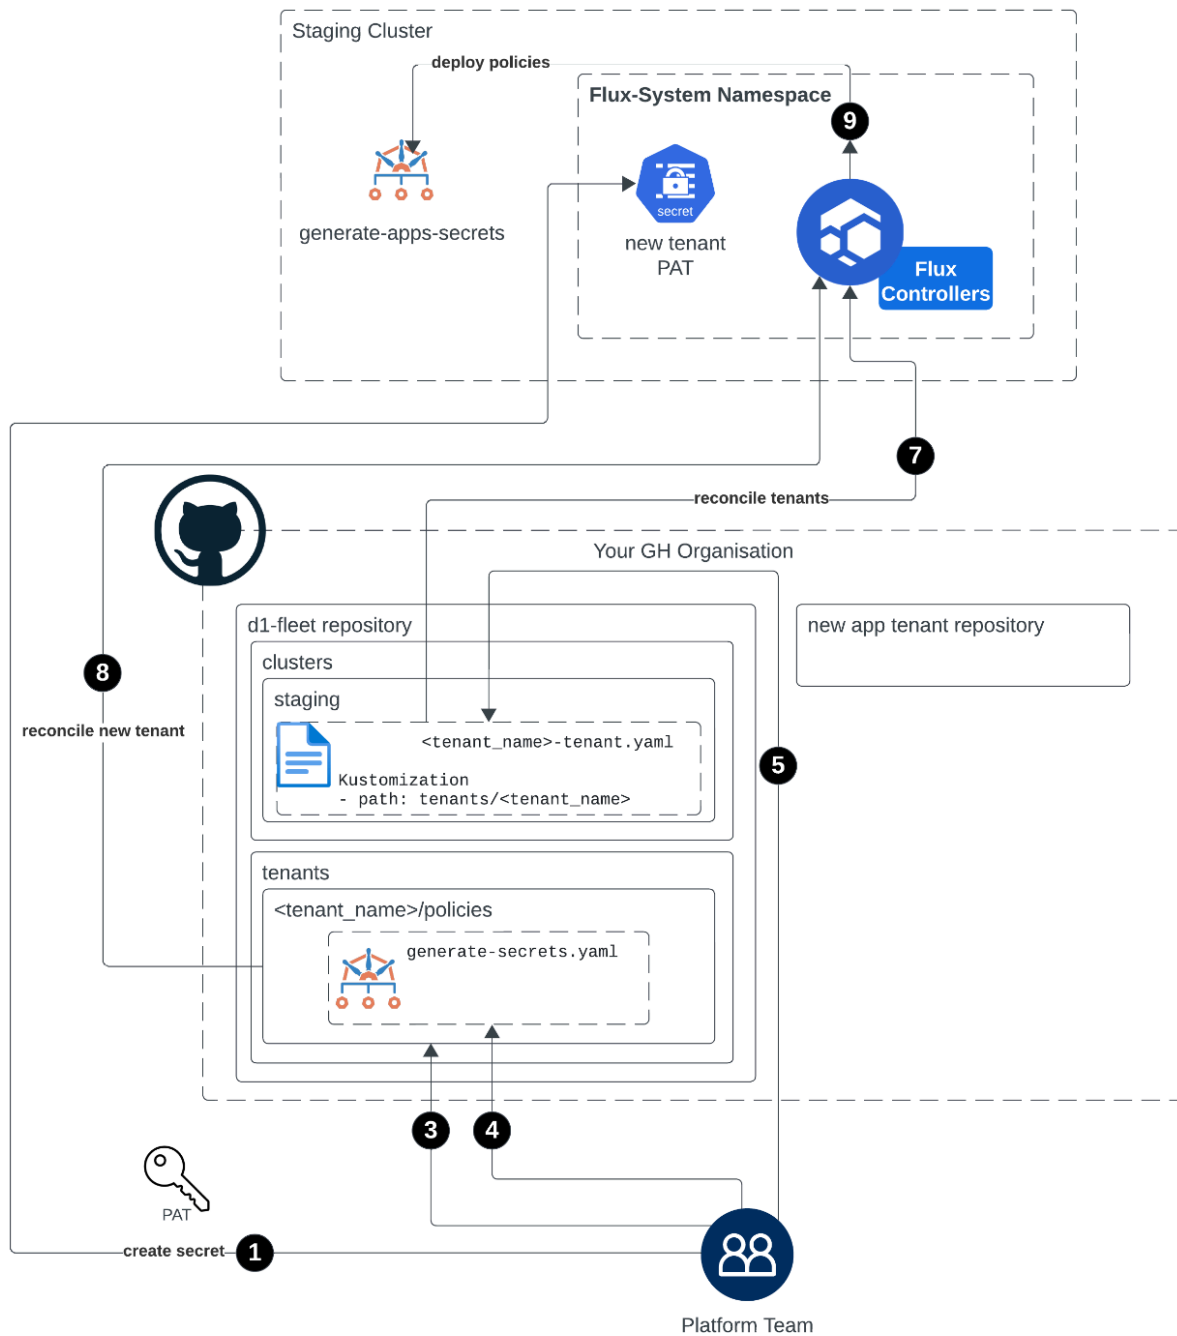


Figure 5: Example Workflow - Adding a new application team/tenant.

To onboard an application tenant the platform should fulfil the following requirement:

- Read-write access to d1-fleet repository
- Knows the PAT for the new tenant (see [here](#) for further considerations)

The platform team should prepare the d1-fleet repo for the new tenant in the following way:

1. Create a secret storing the PAT for this tenant repo on all the affected clusters, via the `flux create secret git` command (see [here](#) for additional consideration on this topic)
2. Create a new folder `[d1-fleet]/tenants/<new_tenant>` that will store all the tenant manifests (i.e. `[d1-fleet]/tenants/apps`. This will be populated by resources for each application (**Namespace**, **GitRepository**, **Kustomization**, **RBAC**, etc)
3. Make sure all the resources deployed by this tenant can be easily filtered by a label
4. Push a Kyverno policy in `[d1-fleet]/tenants/<tenant>/policies` to sync all the secrets stored in flux-system namespace belonging to this tenant to the application namespaces. Doing so will allow Flux controllers to be able to retrieve the needed secrets without crossing namespaces. This can be done using the previously declared label
5. Push a new **Kustomization** CR to each affected cluster in `[d1-fleet]/clusters/<cluster>/<tenant_name>-tenant.yaml` that will act as a pointer to the newly created folders and files
6. (optional) Push the manifests needed for the update automation:
  - a. Dedicated **Namespace apps-update**
  - b. **GitRepository** CR that will act as a pointer to the new tenant Git repository, with proper push permissions on the main branch
  - c. **ImageUpdateAutomation** CR that will update the manifests in the new tenant Git repository in the **components** folder
  - d. **RBAC** (**ServiceAccount** + **RoleBinding**) used to reconcile the update components for this new tenant
  - e. **Kustomization** CR that will start reconciliation for the **update** folder
  - f. Kustomize overlay to include all the resources declared before and to add the `toolkit.fluxcd.io/role: automation` label

Flux will then:

7. Reconcile the tenants and then notice a new **Kustomization** for the new tenant

8. Start reconciling the new tenant folder [d1-fleet]/tenants/<new\_tenant>
9. Reconcile the **ClusterPolicy** CRs

Namespaces and RBAC manifests for the tenant might be created using the [flux create tenant](#) command.

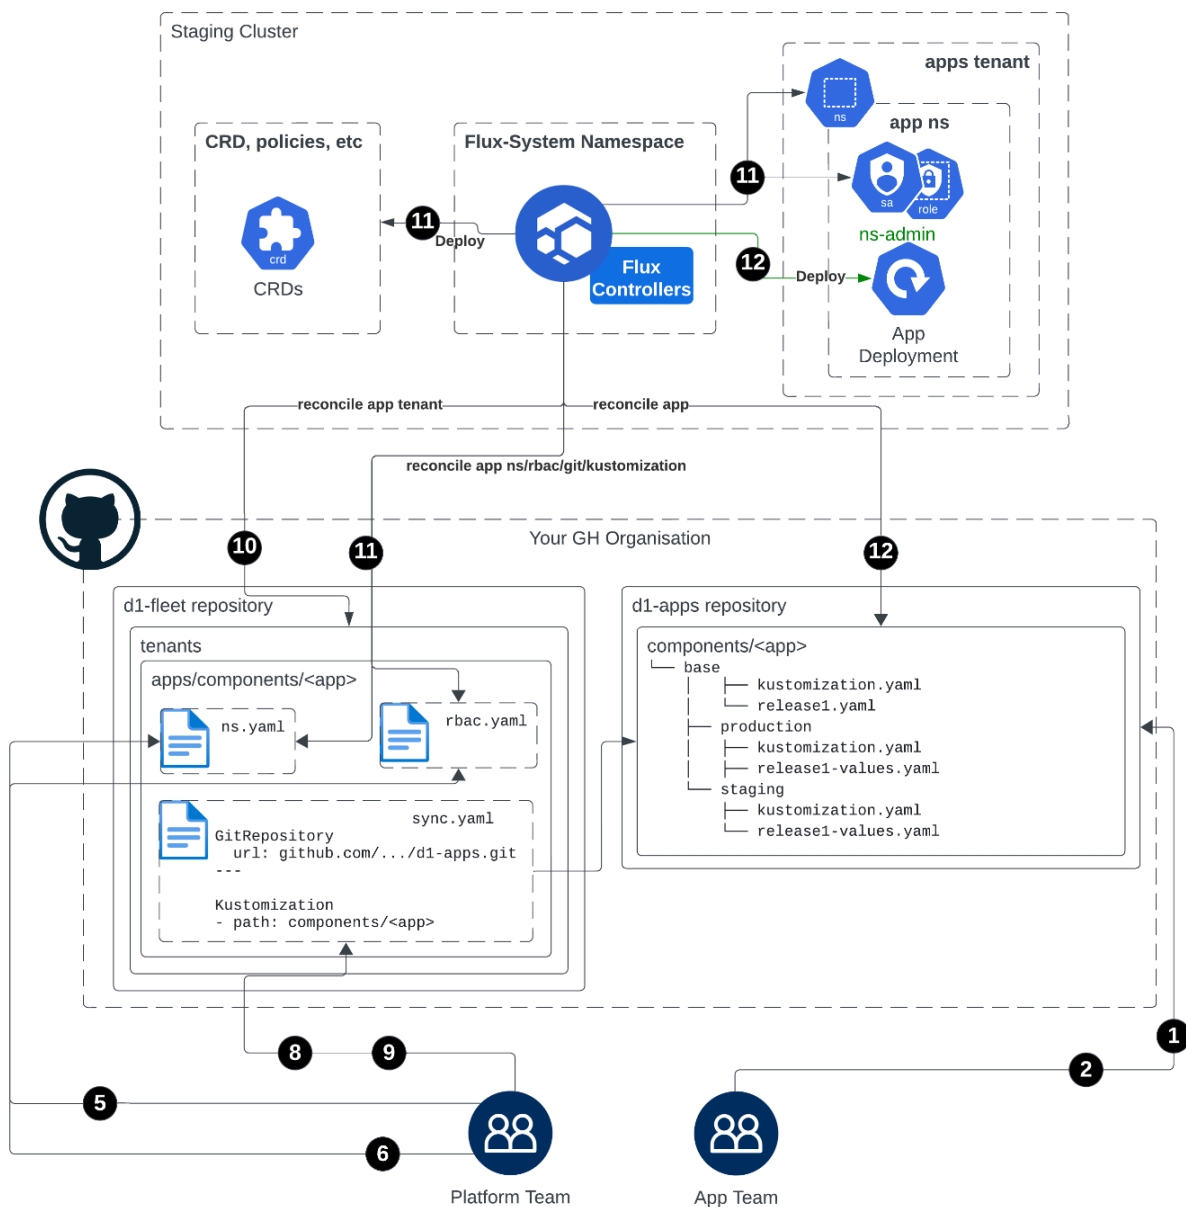
The application team responsible for the new application tenant will have to manage their manifests inside the newly pointed Git repository.

An application tenant might be organised in namespaces in the following way:

```
Unset
.
├── components
│   ├── <namespace>
│   │   ├── base
│   │   ├── production
│   │   └── staging
│   └── <namespace>
│       ├── base
│       ├── production
│       └── staging
└── update
```

## Managing an Application

Manifests that deploy namespaced objects belonging to the app are managed in the proper app tenant. Adding a new app will require a joint work between the platform team and the application team responsible for that application. The application team might have to share with the platform team all the needed resources that can't be deployed by a namespaced admin role and thus cannot be automatically reconciled by the Flux role that will be assigned to the application namespace.



**Figure 6: Example Workflow - Managing an application.**

To onboard a new application in an existing tenant the platform team should fulfil the following requirement:

- Read-write access to d1-fleet repository

The application team should fulfil the following requirements:

- Read-write access to the application tenant

The application team should push on the application tenant:

1. Push the needed manifests (i.e. [HelmRepository](#) and [HelmRelease](#)) on their

- repository
- 2. Push any needed environment overlays
- 3. Pull request on main branch to apply changes on production clusters
- 4. (optional) Push the manifests needed for the update automation (**ImagePolicy** and **ImageRepository**)

The platform team should then push on d1-fleet repo all the manifests to prepare the environment creating non namespaced objects:

- 5. **Namespace**;
- 6. RBAC (**ServiceAccount** + **RoleBinding**) for Flux reconciliation for that specific application
- 7. **GitRepository** CR that will act as a pointer to the tenant Git repository
- 8. **Kustomization** CR that will start the reconciliation for that repo on the application folder
- 9. Kustomize overlay to include all the resources declared before

Make sure all the resources above are tagged with a tenant identifier (see [here](#))

Flux will then:

- 10. Reconcile the application tenant and then notice a new **Kustomization** for the application
- 11. Reconcile RBAC, CRDs and **Namespace**
- 12. Reconcile the objects from [d1-apps]/components/<app> folder (**HelmRepository**, **HelmRelease**, etc)

This would be a possible structure for an application folder in d1-fleet repo:

```
Unset
./tenants/apps/components/<namespace>/
├─ kustomization.yaml
├─ namespace.yaml
├─ rbac.yaml
└─ sync.yaml
```

# Further Implementation Guidance

## Securing GitHub

The D1 reference architecture uses GitHub as the version control system. In order to offer more complete information, the next sections include the main security aspects to be taken into account for a safer Flux integration. To find out more detailed information about all the security controls and how to implement them please visit the [GitHub official documentation](#).

Having permissions to push configuration into the infra or fleet repositories is equivalent to having cluster admin permissions over the cluster, as the account which is reconciling the state has been granted cluster admin permissions. With the d1-apps repository the blast radius would be reduced, as anything the applications team pushes to the d1-apps repository will be reconciled with namespace admin permissions. Hence, securing GitHub becomes critically important within Flux d1, or any GitOps architecture.

It is worth mentioning that the security measures detailed below are an overall guidance but there could be additional considerations to be taken into account for each specific scenario.

## GitHub as a Production Service

It is necessary to analyse the tolerance of the target organisation to a potential Flux outage. A GitHub outage would prevent the deployment of workloads within the onboarded clusters. The fallback would be to resort to manual deployment with administration of the cluster and management of applications conducted by the platform and application teams respectively.

Therefore, analysing the service level agreements offered by GitHub in its different plans is a key factor to take into account within the whole Flux D1 architecture deployment.

## User Access Control and Authorisation

In order to minimise the risk of malicious deployments, any interaction with Flux integrated repositories should only be allowed to those individuals and processes which actually need it. In order to guarantee a proper access control to the Flux repositories the next GitHub controls should be enforced:

- Two-factor authentication: enforce two factor authentication for all the users to prevent unauthorised access to your GitHub repositories
- Team access controls to the repositories: To address proper segregation of duties and observe the least privilege principle, a more fine grained access control is convenient to regulate access to the Flux repositories. Using Teams and then adding the necessary users within the platform and applications allows granting them only the necessary permissions over the desired repositories
  - Limit access to d1-fleet and d1-infra repositories: access to these repositories should be only reserved to platform team members. Ideally, this GitHub team might be used to authorise direct administrative cluster access following a break the glass process, in order to guarantee a 1 to 1 mapping between who has access to these repositories (that will be reconciled with cluster-admin roles) and who can get direct administrative cluster access (via kubectl or consoles)
  - Limit access to application repositories: access to application repositories (that is, repositories that store the manifests that will deploy the applications, i.e. d1-apps) should be reserved to application team(s). Multiple teams will yield to multiple repositories that should be added as separated tenants

In Github, personal accounts have no possibility to create specific roles but the other tiers permit the creation of up to three custom roles to allow more fine grained access control to the Flux repositories. The next is a recommendation of the possible permissions to be granted for the repositories and teams proposed above:

- Applications Team
  - D1-fleet: no permissions
  - D1-infra: no permissions
  - D1-apps: Members with write permissions. At least two admins within the repository or two owners
- Platform Team
  - D1-fleet: write permissions, at least two admins within the repository or two owners
  - D1-infra: write permissions, at least two admins within the repository or two owners
  - D1-apps: see [here](#) for different alternatives

This permissions configuration is considered to meet the least privilege principle but may not be fully suitable for every organisation. If additional permissions are needed



within the target organisation, it is always advisable to keep the access permissions to the minimum possible for both the platform and application team.

If it is necessary to segregate the permissions among different application teams, the best way to do it is by adding a different tenant.

## Protected Branches

Considering the specific Flux integration with GitHub repositories, the next controls are recommended at a branch level:

- D1-fleet Repo → main branch: Pull Requests to be approved by at least two approvers as Flux reconciles from this branch and it is really critical from a service account permissions point of view
- D1-infra Repo → main branch: Pull Requests to be approved by at least one approver. This has the objective of having an extra review in case some manifest could be overly permissive and detect earlier potential deviations. Flux bot has to be able to bypass this rule in order to automate updates on this branch. Further explanation on how to protect the branch and allow the bypass is given at the end of this section
- D1-infra Repo → production branch: Pull Requests to be approved by at least two approvers as this is the branch Flux reconciles the status and it will use cluster admin permissions, in the same way as the fleet repository
- D1-apps Repo → main branch: Pull Requests to be approved by at least one approver, to guarantee there is at least a peer review for any new artefact to be created within the cluster. Flux bot has to be able to bypass this rule in order to automate updates on this branch. Further explanation on how to protect the branch and allow the bypass is given at the end of this section
- D1-apps Repo → production branch: Pull Requests to be approved by at least two approvers. Similarly to what happens with the Flux infra repository, this is the branch the source controller will reconcile from and then deploy the new artefacts using an account with namespace administration permissions. Even though this service account does not have as wide permissions as the one used to reconcile the Infra or Fleet repositories, it is still possible to get control over cluster nodes deploying malicious artefacts and try to escalate privileges from there

In order to protect the branches previously mentioned, it is necessary to create a branch protection rule as next:

1. Go to the corresponding repository within GitHub
2. Click on Settings → branches → new protection rule

3. Check *require a pull request before merging*
  - a. Inside this rule, enable *require approvals*
  - b. Select the number of approvals suggested previously or any other suitable for your organisation
  - c. Select *Allow specified actors to bypass required pull requests*
  - d. Search for the flux-bot-account you created for the bootstrap and add it
4. Finally, populate the *Branch Name Pattern* at the top of the page with any pattern matching the needed criteria, in the previously exposed cases, you can populate with *main*
5. At the bottom of the page click on the *create* button

This will request every user who wants to make a change to the branch to require a pull request which also needs to be approved, but it will also allow the flux-bot-account to bypass the rule and push from the Flux deployed cluster directly to main, using the provisioned PAT.

## Signed Commits

Apart from guaranteeing the minimum privilege principle and the separation of duties, it is also desirable to be able to track who made each change within the files in each repository.

When you work locally on your computer, Git allows you to set the author of your changes and the identity of the committer. This, potentially, makes it difficult for other people to be confident that commits and tags you create were actually created by you. To help solve this problem you can [sign your commits](#) and tags. GitHub marks signed commits and tags with a verification status.

To give other users increased confidence in the identity attributed to each user's commits and tags, it is recommended that all the organisation's users are mandated to enable the *vigilant mode*. With *vigilant mode* enabled, all of the commits and tags are marked with their verification status.

## Flux - Github Role Based Access Control (RBAC)

### Bot Account

As mentioned in the [Cluster Onboarding section](#), a new GitHub account for the Flux bot is required. This account will be used by the Flux CLI and the Flux controllers running on clusters to authenticate with GitHub during cluster bootstrap, fleet (d1-fleet repo) and infrastructure (d1-infra repo) reconciliation.

For convenience, the newly created Flux bot account will be managed by the platform team in your organisation with the following permissions:

- Read and write access to the *d1-fleet* repository (required for cluster bootstrap)
- Push access to the *main* branch of the *d1-fleet* repository (required for cluster bootstrap)
- Read and write access to the *d1-infra* and *d1-apps* repositories (required for cluster and application reconciliation and [image automation](#))

For these permissions to be granted, the bot account must be part of the Github Organisation and also part of an Organisation Team that has read and write permissions to all relevant repositories.

### Personal Access Tokens

To allow Flux to use the bot account, Personal Access Tokens (PATs) must be generated, these are used by the Flux controllers to authenticate using the bot account. To prevent Application tenant privilege escalation through use of a shared PAT, multiple PATs are generated, a single PAT for access to fleet and infra repositories, and a unique PAT for each new application team with access limited to that application team's repositories only.

The platform PAT has the following permissions for the *d1-infra* and *d1-fleet* repositories:

- *Administration* -> Access: Read-only
- *Commit statuses* -> Access: Read and write
- *Contents* -> Access: Read and write
- *Metadata* -> Access: Read-only

This token will be stored in all clusters to authenticate with GitHub to pull the fleet desired state from the *d1-fleet* and *d1-infra* repositories. The token is also used to automate the Helm chart updates in the *d1-infra* repository, where the bot account has push access to the main branch.

To start reconciling the cluster from an application tenant, a new PAT has to be created with the following permissions:

- *Administration* -> Access: Read-only
- *Commit statuses* -> Access: Read and write
- *Contents* -> Access: Read and write
- *Metadata* -> Access: Read-only

but limiting its access to the application repository only (i.e. *d1-apps*).

This token will be stored in all clusters before onboarding the apps tenant (see [here](#)) and will be used by Flux to authenticate with GitHub and to start reconciling the tenant from the *d1-apps* repository. The token is also used to automate the Helm chart updates in the *d1-apps* repository, where the bot account has push access to the main branch (see [here](#) for additional considerations on read-write PAT).

As soon as the clusters need to be reconciled from a new tenant, a new dedicated PAT should be created, with limited access to the repositories belonging to that tenant.

N.B Deploy Keys are an option supported by Flux, but are not considered within this guide owing to a lack of key expiry.

### PAT Management Considerations

The above model is based for convenience on several trade-offs that must be carefully considered:

- PATs are generated by the same bot account that belongs to the platform team with read-write permissions on all the repos. As an alternative that follows the least privilege principle, PATs might be generated from multiple accounts belonging to different teams with limited privileges only on the repositories the PAT will have access to. However, more granularity might require additional effort and less velocity
- PATs will be added as Kubernetes secrets by the platform team during cluster bootstrap and during tenant onboarding. This means that the platform team has to know these tokens. The platform team might be the same team that manages the GitHub organisation, therefore they might be responsible for creating repositories, teams, bot accounts and the PATs too. The team will then inherit read and write access to all the repositories, regardless of the tenant. This simplified model might not fit well with all the organisations, if a stricter segregation of duties is in place
- Mixed and more complex responsibility models affect how, when and by whom repositories/accounts and PATs are created. As an example, credentials and/or PATs might be created by a dedicated GitHub team with administrative access to the organisation but with no direct access to the clusters. Therefore, secrets - might them be bot credentials or ready-to-use PATs - will have to be securely shared between the GitHub and the platform team, affecting how clusters are bootstrapped and/or tenants are onboarded
- As PATs expire by design, when deciding on a PAT management model, rotation must be taken into account. For a single platform team managed bot

account, it should be possible for PATs to be rotated en masse, assisted by automation. Should separate teams be responsible for different bots and their respective PATs, then rotation will be their responsibility. No artefact will be pruned by Flux if the token is not renewed before the expiration date. However, no further update from the repo will be reconciled by the cluster until the PAT is renewed.

## Extending to other Git Providers

This model is specifically intended to work with GitHub permission models. This might be adapted to other Git providers as long as the PAT and team permissions are kept as similar as possible and follow the least privilege principle. Any consideration of SLA for the alternative provider is left to the enterprise.

## Multi-Tenancy Considerations

### Multi-Tenancy Configuration of Flux

#### Kustomization

The D1 model relies on the [Flux multi-tenancy configuration](#). As far as this model is concerned the tenants have restricted access to the cluster(s) according to the Kubernetes RBAC configured by the platform admins. The repositories owned by tenants are reconciled on the cluster(s) by Flux, under the Kubernetes account(s) assigned by platform admins.

According to the official documentation, the Flux installation is locked down using the following patches during the bootstrap process (see [here](#) for production clusters Kustomize overlay):

- Deny cross-namespace access to Flux custom resources, thus ensuring that a tenant can't use another tenant's sources or subscribe to their events.
- Deny accesses to Kustomize remote bases, thus ensuring all resources refer to local files, meaning only the Flux Sources can affect the cluster-state.
- All **Kustomizations** and **HelmReleases** which do not have `spec.serviceAccountName` specified, will use the default account from the tenant's namespace for deployment. Tenants have to specify a service account in their Flux resources to be able to deploy workloads in their namespaces as the default account has no permissions.
- The **flux-system Kustomization** is set to reconcile under a service account with **cluster-admin** role, allowing platform admins to configure cluster-wide resources and provision the tenant's namespaces, service accounts and RBAC.

## Default Service Account

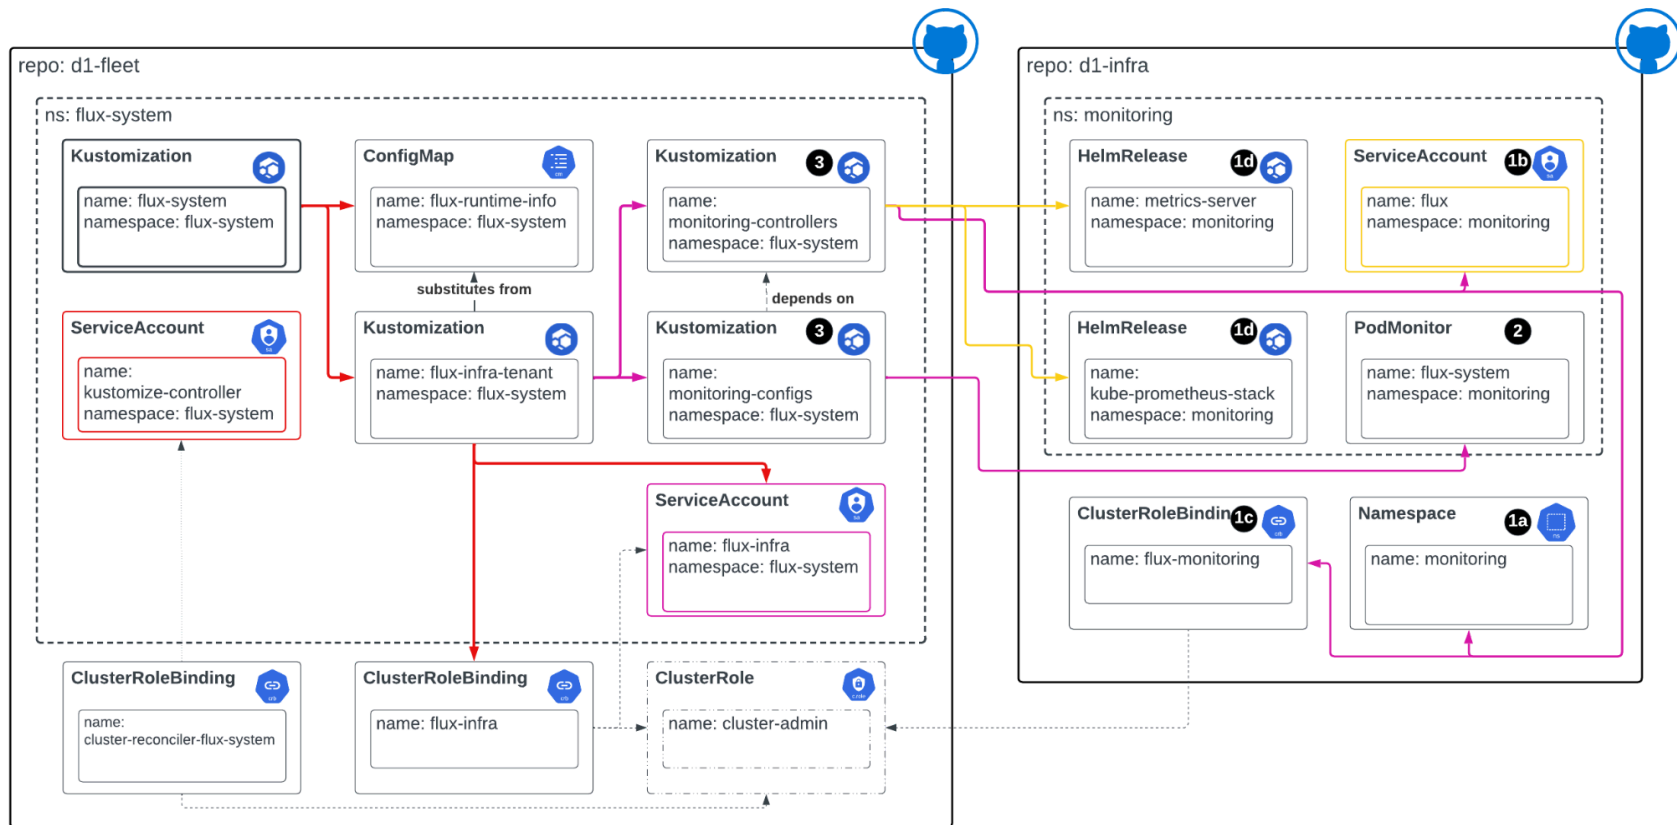
Even if application tenants are reconciled with service accounts with namespaced admin roles, particular attention should be always paid to the `default` service account. `default` is the fallback account for `Kustomization` and `HelmReleases` with no `.spec.serviceAccountName` specified, therefore this should be carefully monitored to guarantee no unnecessary roles are assigned to it.

## Flux Impersonation Model


In d1, Flux uses multiple service accounts to act on the cluster:

- `default` in `flux-system` is used as fallback service account when no service account is explicitly specified in `spec.serviceAccountName`
- `kustomize-controller` in `flux-system` namespace is used to reconcile `d1-fleet`, sync Flux components and to reconcile the tenant pointers. In addition, this is used to create non namespaced resources for the application tenant. This service account is used to reconcile the cluster policy that will sync needed secrets and configmaps across the namespaces
- `flux-infra` in each infrastructure namespace is used to reconcile the infrastructural controllers (details in Figure 6)
- `flux-apps` in each application namespace is used to reconcile the applicative components (details in Figure 7)

The multi-tenancy lock will allow Flux to only impersonate service accounts belonging to the same namespace where the `Kustomization/HelmRelease` is being deployed to.



**Figure 7: Flux impersonation model - flux-infra.**



**Figure 8: Flux impersonation model - flux-apps.**



## Flux Controller Kubernetes RBAC Considerations for Multi-Tenancy

As part of onboarding an infrastructure component, or a namespace for an application team to the cluster, the platform team will define the permissions and service account used by the kustomize-controllers and helm-controllers to perform the deployments within Kubernetes. These are defined within the [dl-fleet\\_repo](#) (example in link)

Within the provided example repos, it is important to note that these are namespace admin permissions for application tenants and cluster admin permissions for infrastructure components. It is recommended that you analyse these permissions and define your own roles based upon your use case for the following reasons:

- These permissions are not configured in line with least-privilege
- If Infrastructure components are managed by external teams, providing those teams with cluster admin permissions via flux, provides an escalation route for non platform team members
- `admin ClusterRole` can vary from distribution to distribution. Binding to role that includes `*`'s is strongly discouraged
- It is likely that both application and platform teams will need to configure their applications or infrastructure components to use custom resources such as cert-manager CRs, and these will be needed to be added to the roles that Flux impersonates so that deployment can occur

## Policy Pack for Multi-Tenancy

The following Kyverno policies are stored within the *dl-fleet/tenants/apps/policies*. Each one of them has a specific purpose and might require additional consideration:

### Secret Cloning and Least Privilege

Application tenants are able to access the Git repository via a secret stored in the application (i.e. backend/frontend) namespaces. The D1 model currently requires the platform team to create these secrets (via `flux create secret git` command) in the flux-system namespace and then let a Kyverno policy sync them in the applicative namespaces belonging to that tenant ([here](#) for more details). The secret stores a single PAT with read-write permissions per tenant, for both reconciliation (fetch updates from the repo and apply them to the cluster) and update automation (update manifests on the main branch).

An attacker might be able to exploit a vulnerable application with enough privileges to access this read-write PAT and use it to affect the tenant main branch, directly affecting all the staging clusters. As an alternative, an attacker might be able to

mount the secret in an arbitrary application, if no guardrails are implemented (i.e. admission controller policy).

It would be preferable to use a read only PAT for reconciliation - synced to all the application namespace belonging to that tenant - and a read/write PAT for the update automation workflow - stored in a single namespace where no workflow is executed.

## Runtime Information

`flux-runtime-info` ConfigMap can be used by Kustomization and therefore by workloads (infra or applicative). DI does not allow Flux controllers to make cross-namespace references, therefore, similarly to what happens for PAT secrets, the ConfigMap is automatically copied in all the namespaces by Kyverno policy (`sync-flux-configmaps`), making the content of the ConfigMap potentially available to every workloads.

Therefore no confidential information should be stored in this ConfigMap, in line with best practice

## Other Security and Availability Considerations

### Flux CLI Usage & Alternatives

The guide describes cluster bootstrapping using the Flux CLI from a platform team members device, with the Flux CLI consuming the `kubeconfig` file on the device as well as the Github Flux Bot Personal Access Token.

With phishing and administrator device compromise becoming an increasingly common attack vector, devices should be hardened appropriately for this use case, with the platform team subject to security awareness training.

Within regulated organisations, handling of privileged Kubernetes credentials, such as cluster-admin for production clusters, would not normally be permissible on a platform team members device. Thus use of the Flux CLI on a platform team members laptop consuming privileged credentials for normal operations such as cluster bootstrap, would not be permissible. Alternatives include using a dedicated Bastion VM in the same VPC as the cluster to run the CLI or the [Flux Terraform provider](#), or using CI runners with credentials (PAT and Kubernetes credentials) stored as CI secrets.

## Update Automation Service Level

In the D1 architecture, the controllers responsible for the update automation workflow are only deployed in staging clusters (see [here](#)). That means that the manifests stored on the repo are automatically going to be updated only if at least one staging cluster is available. This is something that should be taken in consideration and that would depend on which kind of service availability is expected for the update automation process.

## Tenant Denial of Service Abuse Case

Given that Flux controllers will automatically reconcile a cluster with the contents of a source git repository there is a situation that exists where should a repository be arbitrarily filled with a large number of Kubernetes resources, this could prevent other tenant deployments to the cluster through increased Flux processing and DOS of the Kubernetes API endpoint.

This would occur regardless of any Kubernetes admission control policies as Flux controllers would need to pull, template the resources, authenticate and send to the K8s API, which would then receive the resources and parse them before applying any mutating or validating policies.

In order to mitigate this potential problem, it is considered a good practice to set quotas for the total number of certain resources of all standard, namespaced resource types within the application namespaces by using an [object count quota](#).

[API Priority and Fairness \(APF\)](#) is another feature which mitigates a potential Denial of Service against the API Server. APF introduces a limited amount of queuing, so that no requests are rejected in cases of very brief bursts. Requests are dispatched from queues using a fair queuing technique so that, for example, a poorly-behaved controller need not starve others (even at the same priority level).

# About ControlPlane Flux Enterprise

January 2024, at ControlPlane we launched ControlPlane Enterprise for Flux CD (CPE-Flux CD), designed to enhance highly regulated organisations who run and manage production Kubernetes environments. CPE-Flux CD provides a secure, stable, and efficient subscription based offering for enterprise Kubernetes environments, while ensuring the long term sustainability of the upstream project.

We provide this by employing core maintainers of the project and offering additional security and vulnerability remediation enhancements. If your organisation requires or benefits with these enhancements over the upstream Flux CD project, CPE-Flux is an option to consider:

- Enhanced security: Security guardrails, hardened and distroless container images, and FIPS-compliant Flux CD builds.
- Vulnerability management: Continuous scanning for CVEs, patching of Flux CD, and SLAs for critical vulnerabilities with signed VEX documentation.
- Seamless upgrades and maintenance: Zero-downtime migration, full upstream compatibility, OCI-compliant image repositories, compatibility with the latest six Kubernetes releases, and SLAs for critical vulnerability remediation.
- Dedicated support: Developer Support Portal for tracking and managing resolution and feature requests.

All this additional support and enhanced security features makes Flux Enterprise a great choice for those companies aiming for the operational excellence of their Kubernetes environments, specially for those which are dealing with demanding environments such as:

- Security critical workloads requiring a clear segregation of duties
- Necessity of enforcing the least privilege principle at Service Account Level
- Large environments with at scale deployments
- CI/CD pipelines with high speed delivery
- Lack of highly specialised personnel capable to debug low level Flux open source potential integration issues
- New featured deployments are availability critical
- Standards compliance or regulations mandate any third party integration have to be compliant with the best security practices

Additionally, CPE-Flux CD is backed by ControlPlane's Technical Account Architect (TAA) offering and professional services team, providing additional resources and support as needed.

It is important to emphasise that while CPE-Flux CD offers these additional security and remediation enhancements, that if they are not a requirement in your organisation you should use the upstream Flux CD project.

Keep in mind that ControlPlane's team is available to support you in achieving secure, stable, and efficient Kubernetes production environments, whether you choose to use CPE-Flux or the upstream Flux CD project. By offering CPE-Flux, ControlPlane aims to ensure the sustainability and continued development of the upstream Flux CD project.

# Appendix

## Administrative Tasks

### Bootstrapping a Staging Cluster

Make sure to set the default context in your `kubeconfig` to your staging cluster, then run bootstrap with:

```
Unset
export GITHUB_TOKEN=<Flux Bot PAT>

flux bootstrap github \
  --registry=ghcr.io/fluxcd \
  --components-extra=image-reflector-controller,image-automation-controller \
  --owner=<owner> \
  --repository=d1-fleet \
  --branch=main \
  --token-auth \
  --path=clusters/staging
```

This command will explicitly enable `image-reflector-controller` and `image-automation-controller`. See [here](#) for more details.

### Bootstrapping a Production Cluster

Make sure to set the default context in your `kubeconfig` to the production cluster you want to install Flux to, then run bootstrap with:

```
Unset
export GITHUB_TOKEN=<Flux platform PAT>

flux bootstrap github \
  --registry=ghcr.io/fluxcd \
  --owner=<owner> \
  --repository=d1-fleet \
  --branch=main \
  --token-auth \
  --path=clusters/prod-eu
```

## Rotating the PATs

To rotate the main flux-system secret (used to connect to the fleet and infra repository) on the currently selected cluster in [kubecfg](#), the platform team should run this command:

```
Unset
flux create secret git flux-system \
  --namespace=flux-system \
  --url=https://github.com \
  --username=git \
  --password=$NEW_GITHUB_TOKEN
```

To rotate the secrets used by Flux to connect to the other tenant's repository on the currently selected cluster in [kubecfg](#), the platform team should run this command:

```
Unset
flux create secret git flux-<tenant> \
  --namespace=flux-system \
  --url=https://github.com \
  --username=git \
  --password=$NEW_GITHUB_TOKEN
```

These commands should be executed for each of the affected clusters.

The [Kyverno policies](#) deployed by D1 will make sure that the secrets are propagated in the needed namespaces.

## Adding an Infra Component

In this example, we are going to deploy MetalLB using Flux. [MetalLB](#) is *is a load-balancer implementation for bare metal [Kubernetes](#) clusters, using standard routing protocols.*

To do so, the platform team should perform the following actions on fleet and infra repositories:

### Infra repository

- Add [namespace.yaml](#) in d1-infra/components/metallb/base. This will contain:
  - metallb [Namespace](#) definition
  - [ServiceAccount](#) and [ClusterRoleBinding](#) for Flux reconciliation

```

Unset
apiVersion: v1
kind: Namespace
metadata:
  name: metallb
  labels:
    app.kubernetes.io/component: lb
    toolkit.fluxcd.io/tenant: platform-team
    pod-security.kubernetes.io/enforce: baseline
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    app.kubernetes.io/component: lb
    toolkit.fluxcd.io/tenant: platform-team
  name: flux-lb
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: flux
  namespace: metallb
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: flux
  namespace: metallb
  labels:
    app.kubernetes.io/component: lb
    toolkit.fluxcd.io/tenant: platform-team

```

- Add `metallb.yaml` to `d1-infra/components/metallb/base`. This will contain `HelmRepository` and `HelmRelease`, with the common values used by all the environments

```

Unset
apiVersion: source.toolkit.fluxcd.io/v1beta2
kind: HelmRepository
metadata:
  name: metallb

```



```

    namespace: metallb
spec:
  interval: 12h
  url: https://metallb.github.io/metallb
---
apiVersion: helm.toolkit.fluxcd.io/v2beta2
kind: HelmRelease
metadata:
  name: metallb
  namespace: metallb
spec:
  serviceAccountName: flux
  interval: 1h
  chart:
    spec:
      version: "0.14.3"
      chart: metallb
    reconcileStrategy: ChartVersion
    sourceRef:
      kind: HelmRepository
      name: metallb
    interval: 12h
  install:
    crds: Create
    timeout: 9m
  upgrade:
    crds: CreateReplace
    timeout: 9m

```

- Add the IP pool for MetalLB as a cluster variable in runtime-info.yaml

```

Unset
apiVersion: v1
kind: ConfigMap
metadata:
  name: flux-runtime-info
  namespace: flux-system
  labels:
    toolkit.fluxcd.io/runtime: "true"
  annotations:
    kustomize.toolkit.fluxcd.io/ssa: "Merge"
data:
  ENVIRONMENT: "staging"
  GIT_BRANCH: "main"

```

```
CLUSTER_NAME: "staging-1"
CLUSTER_DOMAIN: "preview1.example.com"
LB_IP_POOL: "192.168.10.0/24"
```

- Add the IPAddressPool and BGPAdvertisement CRs, in d1-infra/components/metallb/configs

```
Unset
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: main-pool
  namespace: metallb
spec:
  addresses:
  - ${LB_IP_POOL}
---
apiVersion: metallb.io/v1beta1
kind: BGPAdvertisement
metadata:
  name: bgpadvertisement-basic
  namespace: metallb
spec:
  ipAddressPools:
  - main-pool
```

- Add any needed environment overlays in **production** and **staging** folder, on both **controllers** and **configs**

## Fleet Repository

- Add new **Kustomizations** /tenants/infra/components/lb.yaml

```
Unset
tenants/infra
├── components
│   ├── admission.yaml
│   ├── kustomization.yaml
│   ├── lb.yaml # <--
│   ├── monitoring.yaml
│   ├── rbac.yaml
│   └── source.yaml
└── update
```

```
|— automation.yaml
|— kustomization.yaml
└— sync.yaml
```

Unset

```
---
apiVersion: kustomize.toolkit.fluxcd.io/v1
kind: Kustomization
metadata:
  name: lb-controllers
spec:
  serviceAccountName: flux-infra
  interval: 1h
  retryInterval: 2m
  timeout: 10m
  prune: true
  wait: true
  sourceRef:
    kind: GitRepository
    name: flux-infra
    path: components/metallb/controllers/${ENVIRONMENT}
---
apiVersion: kustomize.toolkit.fluxcd.io/v1
kind: Kustomization
metadata:
  name: lb-configs
spec:
  serviceAccountName: flux-infra
  dependsOn:
    - name: lb-controllers
  interval: 1h
  retryInterval: 2m
  timeout: 5m
  prune: true
  wait: true
  sourceRef:
    kind: GitRepository
    name: flux-infra
    path: components/metallb/configs/${ENVIRONMENT}
```

- Add a reference to `lb.yaml` in `d1-fleet/tenants/infra/components/kustomization.yaml` to kick start the reconciliation

Unset

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
namespace: flux-system
resources:
- rbac.yaml
- source.yaml
- admission.yaml
- monitoring.yaml
- lb.yaml # <--
labels:
- pairs:
    toolkit.fluxcd.io/tenant: infra
```

# About

ControlPlane is a global cloud native and open source cybersecurity consultancy operating in London, New York, and Auckland. We have industry-leading expertise in the architecture, audit, and implementation of zero trust infrastructure for regulated industries. With a deep understanding of secure-by-design and secure-by-default cloud, Kubernetes, and supply chain security we conduct threat modelling, penetration testing, and cloud native security training to the highest standard.

ControlPlane is trusted as the partner of choice in securing: multinational banks; major public clouds; international financial institutions; critical national infrastructure programs; multinational oil and gas companies, healthcare and insurance providers; and global media firms.

## Team

Andrea Martino  
Miguel Ángel Hernández Ruiz  
Rowan Baker

## Reviewers

Andrew Martin  
Eduardo Olarte  
Stefan Prodan  
Martin Stadler